

Anomalies Detection

Advanced Machine Learning and Artificial Intelligence

Yuri Balasanov

University of Chicago, MScA

© Y. Balasanov, 2018

© Yuri Balasanov, 2018

All Rights Reserved

No part of this lecture notes document or any of its contents may be reproduced, copied, modified or adapted without the prior written consent of the author, unless otherwise indicated for stand-alone materials.

The content of these lectures, any comments, opinions and materials are put together by the author especially for the course Linear and Nonlinear Statistical Models, they are sole responsibility of the author, but not of the author's employers or clients.

The author cannot be held responsible for any material damage as a result of use of the materials presented in this document or in this course.

For any inquiries contact the author, Yuri Balasanov, at yuri.balasanov@iLykei.com .

Outline of the Session

- Autoencoders and their applications
- General architecture of autoencoders
- Stacked autoencoders
- Sequential training of stacked autoencoders
- Unsupervised pretraining with stacked autoencoders
- Denoising autoencoders
- Sparse autoencoders
 - Kullback-Leibler Divergence
 - Application of Kullback-Leibler Divergence in sparse autoencoders
- Other types of autoencoders

Main text: Hands-On Machine Learning with Scikit-Learn and TensorFlow: Concepts, Tools, and Technologies to Build Intelligent Systems, Aurelien Geron, 2017

Autoencoders and Their Applications

- Autoencoders are artificial neural networks capable of unsupervised learning of representation of inputs (**codings**)
- Codings typically have lower dimensionality than the input data. This makes them a popular method for reduction of dimensionality
- Autoencoders are also commonly used as feature detectors and feature creators
- Autoencoders can help pretraining lower levels of deep neural networks, especially when data have limited labels
- Autoencoders can be used as **generative models**, i.e. they can generate new data similar to the training data
- Finally, autoencoders have become a popular method for anomalies detection

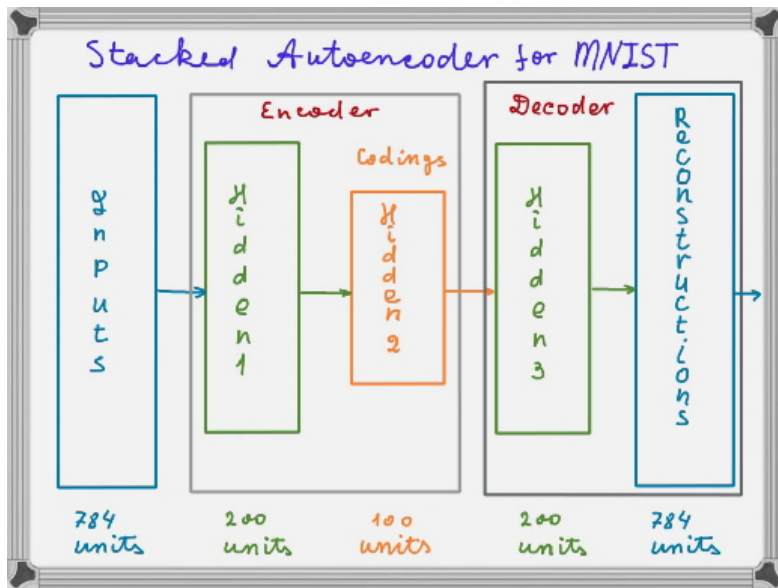
General Architecture of Autoencoders

- Autoencoder contains 2 parts:
 - Encoder (**recognition network**), converts inputs to an internal representation
 - Decoder (**generative network**), converts internal representation to the outputs
- Since autoencoder trains making outputs equal to inputs, numbers of inputs and outputs must be equal
- Outputs of autoencoder are also called **reconstructions** and the loss function is **reconstruction loss**
- Autoencoder is typically **undercomplete**, i.e. internal representation has lower dimensionality than inputs. This prevents trivial solutions when inputs are simply copied to outputs
- Autoencoder can be undercomplete in different ways:
 - Layer of internal representations is smaller than number of inputs (**bottleneck**)
 - Internal representation layer has many units, but has smaller number of connections (**sparse autoencoder**)

Stacked Autoencoders

- Autoencoder that has both encoder and decoder containing multiple layers of neurons is called **stacked autoencoder** or **deep autoencoder**
- Typically stacked autoencoders have equal layer sizes for layers symmetrically located relative to the bottleneck
- In addition weights of decoding layer often equal to transposed weights of encoding layer symmetrical to it. This is called **tying weights**. Tying weights reduces number of weights in the network and speeds up training
- Motivation behind additional layers is the same as for deep neural networks: higher complexity the subject for learning with fewer parameters
- However, adding too many layers may result in trivial solutions or loss of efficiency due to lack of undercompleteness

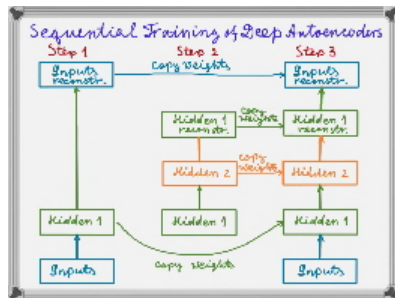
Stacked Autoencoder Example



Sequential Training of Stacked Autoencoders

One additional convenience of training stacked autoencoders is that they can be trained sequentially:

- 1 Train first layer of deep autoencoder as a shallow one, reconstructing the inputs
- 2 Train second shallow autoencoder, reconstructing outputs of the first one
- 3 Stack layers of the deep autoencoder in the right order

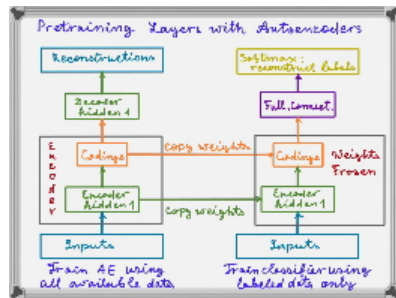


Unsupervised Pretraining with Autoencoders

If there is not enough labeled data use pretrained lower layers by autoencoder:

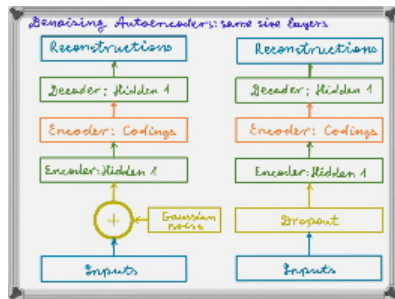
- 1 Use all train data to train an unsupervised autoencoder
- 2 Use the encoder layers as low level layers
- 3 Train higher layers using labeled data

See example for a classification model



Denoising Autoencoders

- Autoencoders can work even if they are not undercomplete
- In order to enforce learning features, but not allowing trivial solution if decoder size is not smaller than size of inputs:
 - 1 Add noise to inputs, or
 - 2 Pass inputs through a dropout layer



Sparse Autoencoders

- Sparsity is another way of putting constraints pushing autoencoders to learn features
- Sparse autoencoder is motivated to reduce the number of active neurons in the coding layer by a type of regularizer added to the cost function
- Regularizer may reduce the number of active neurons in encoder to, say, 5% on average. As a result each input is reproduced as a combination of a limited number of activations
- Let $a_i^{(l)}(x)$ be activation (output) of neuron i in layer l for input x . Sparsity at each training iteration can be measured as average activation in the encoding layer over a batch, assuming the batch size is not too small: $q_i = \frac{1}{m} \sum_{j=1}^m a_i^{(l)}(x_j)$, where x_1, \dots, x_m are inputs of batch of size m
- Given the mean activation per neuron regularizer punishes neurons that are too active by adding sparsity loss to the cost function

Kullback-Leibler Divergence I

- A common sparsity loss function is **Kullback-Leibler divergence (KL-divergence)**. It has much stronger gradient than mean squared error

Definition

Let P and Q be two binomial random variables with probabilities of success p and q correspondingly. Then KL-divergence between them is

$$D_{KL}(P\|Q) = p \ln \frac{p}{q} + (1 - p) \ln \frac{1 - p}{1 - q}$$

- This divergence measure has property that when $p = q$ then $D_{KL}(P\|Q) = 0$ and $D_{KL}(P\|Q)$ monotonically increases very quickly as p and q become more different from each other

Kullback-Leibler Divergence II

Kullback-Leibler Divergence:
Binomial Distributions.

Definition: $\mathcal{D}_{KL}(P||Q) = E_p[\ln \frac{P}{Q}] = \int p(x) \ln \frac{P(x)}{q(x)} dx$

$$P(x) = \binom{n}{x} p^x (1-p)^{n-x} \quad q(x) = \binom{n}{x} q^x (1-q)^{n-x}$$

$$\begin{aligned} \mathcal{D}_{KL}(P||Q) &= \sum_{i=0}^n \binom{n}{i} p^i (1-p)^{n-i} \ln \frac{p^i (1-p)^{n-i}}{q^i (1-q)^{n-i}} \\ &= \sum_{i=0}^n \binom{n}{i} p^i (1-p)^{n-i} \left[i \ln \frac{p}{q} + (n-i) \ln \frac{1-p}{1-q} \right] \\ &= \ln \frac{p}{q} \sum_{i=0}^n i \binom{n}{i} p^i (1-p)^{n-i} + \ln \frac{1-p}{1-q} \sum_{i=0}^n (n-i) \binom{n}{i} p^i (1-p)^{n-i} \\ &= \boxed{\ln \frac{p}{q} np + \ln \frac{1-p}{1-q} n(1-p)} \quad \text{Or for } n=1 \quad \boxed{\ln \frac{p}{q} p + \ln \frac{1-p}{1-q} (1-p)} \end{aligned}$$

$$\begin{aligned} \ln \frac{p^i (1-p)^{n-i}}{q^i (1-q)^{n-i}} \\ = i \ln \frac{p}{q} \\ + (n-i) \ln \frac{1-p}{1-q} \end{aligned}$$

Application of KL-Divergence in Sparse Autoencoders

- To train sparse autoencoder it is necessary to minimize loss equal to sum of KL-divergences between the target activation p (usually a small number) and actual mean activations q_i of neurons i in the encoding layer

$$D_{KL}(P\|Q) = \sum_{i=1}^k \left(p \ln \frac{p}{q_i} + (1 - p_i) \ln \frac{1 - p_i}{1 - q_i} \right)$$

- For sparse autoencoders it is important to have activations $a_i^{(l)}(x)$ for all neurons strictly between 0 and 1. If activation equals either 0 or 1 exactly $D_{KL}(P\|Q)$ may return NaN
- A simple solution ensuring that activations are not equal 0 or 1 is to use logistic sigmoid as activation function for the encoding layer
- A good reference on sparse autoencoders is the following lecture notes article:



Sparse autoencoder, Andrew Ng, CS294A Lecture notes,
<https://web.stanford.edu/class/cs294a/sparseAutoencoder.pdf>

Other Types of Autoencoders

- Development of new variations of autoencoders have been going on adding to their success
- Some of the variations are below:
 - **Contractive autoencoder:** Constrains derivatives of the codings with respect to inputs
 - **Stacked convolutional autoencoder:** Extracts visual features by reconstructing images decoded by convolutional layers
 - **Generative stochastic network:** Denoising autoencoder with generative capability
 - **Winner-Take-All (WTA) autoencoder:** Only $k\%$ activations of each neuron over training batch are kept
 - **Generative Adversarial Network (GAN):** Two networks: "generator" and "discriminator" - play against each other, one trying to present either objects from the real sample or generate fake data; the other trying to detect fakes. Makes a very high quality generative model